REPORT MACHINE LEARNING AND PATTERN RECOGNITION

Fingerprint spoofing detection

Alessia Intini

s309895

Politecnico di Torino

2023/24

Contents

MLPR Project:
Fingerprint spoofing detection

The goal of project is to perform a binary classification on fingerprint spoofing detection, that is to identify genuine vs counterfeit fingerprint images. The dataset consists of labeled samples corresponding to the genuine (True, label 1) class and the fake (False, label 0) class, the data is 6-dimensional.

1 Dataset Analysis

1.1 Training and evaluation sets

The datasets provided contain 6000 samples; the first 6 values in each row represent the features, while the last value is the label. Specifically they are:

- Training Set: 2990 samples beloging to the Fake class (label 0) and 3010 samples belonging to the Genuine class (label 1)

- Evaluation Set: 3010 samples beloging to the Fake class (label 0) and 2990 samples belonging to the Genuine class (label 1)

We will use the Training Set to perform all the analysis, during this phase the dataset is divided to use about 60% of it as the training dataset and the remaining 40% for validation. After this step, the most promising models were chosen and the evaluation dataset was used to evaluate them and make the final considerations.

1.2 Features analysis

We can start to analize all features, here is the graph of each feature, through these histograms it is shown how each feature is distributed:
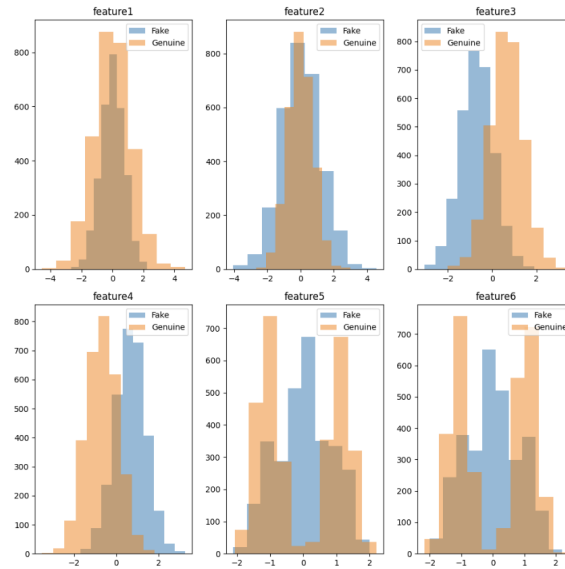

Figure 1: All features

By plotting a histogram for each feature, separately for the Fake and Genuine classes, it is possible

to show whether the samples follow a Gaussian distribution and to what extent. Thus, we can say that some features follow a Gaussian distribution. Furthermore, it can be seen that features 4 and 5 may be the most discriminating features based on their ability to separate the data of the two classes.
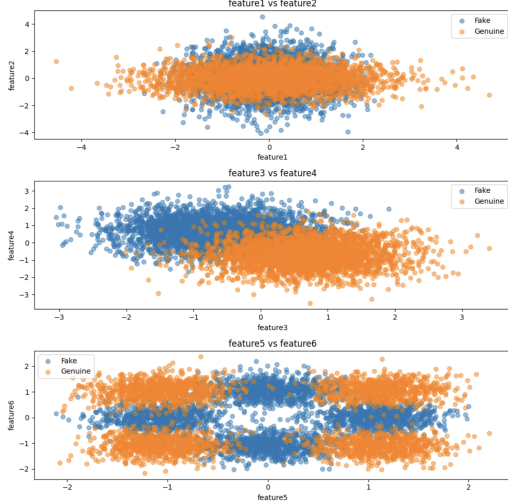

Figure 2: Distribution of feature pairs

Feature	Mean	Variance
1	0.00170711	1.00134304
2	0.00503903	0.9983527
3	-0.00560753	1.0024818
4	0.00109537	0.99029389
5	-0.00700025	1.00119747
6	0.00910515	0.99722374

Figure 3: Mean and variance for each features

After making these calculations we can conclude that the data are already centered since all features have mean close to zero.

Analyzing the pairwise averages and variances of the features, we can make some observations:

- Feature 1 and Feature 2: These two features have averages very close to zero, indicating that their values are similarly distributed around zero. However, the variance of Feature 2 is slightly lower than that of Feature 1, indicating that the values of Feature 2 are slightly less dispersed than those of Feature 1.

- Feature 3 and Feature 4: Feature 3 has a negative mean, while Feature 4 has a positive mean. This might indicate that the values of Feature 3 tend to be lower than those of Feature 4. Also, the variance of Feature 3 is slightly higher than that of Feature 4, indicating that the values of Feature 3 are more dispersed than those of Feature 4.

- Feature 5 and Feature 6: Feature 5 has a negative mean, while Feature 6 has a positive mean. This might indicate that the values of Feature 5 tend to be lower than those of Feature 6. Also, the variance of Feature 5 is slightly higher than that of Feature 6, indicating that the values of Feature 5 are more dispersed than those of Feature 6.

1.3   Features correlation

We can show how correlated the features are by using a heatmap that shows a darker color within the cells for which there is a high correlation between feature i and feature j. The correlation

between feature x and feature y is calculated using Pearson's coefficient:

$$Corr_{i,j} = \frac{Cov(i,j)}{\sqrt{Var(i)}\sqrt{Var(j)}} \tag{1}$$
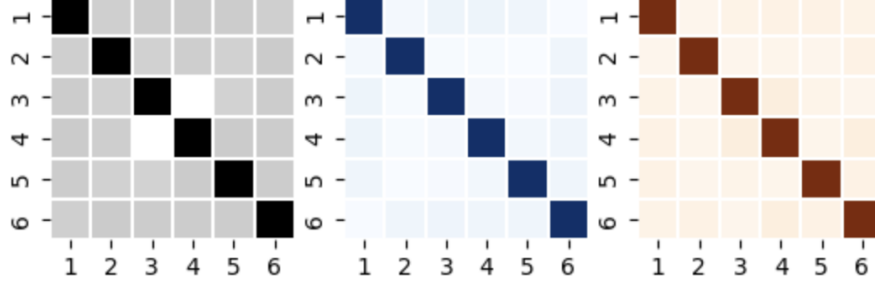

Figure 4: Correlation features: gray scale indicates the correlation between whole dataset, blue is for Fake class and orange for Genuine class

Most of the features do not appear to be particularly correlated; in fact, most of the values in the correlation matrix are close to zero.

Since the features do not appear to be strongly correlated since precisely the values are close to zero. This suggests to us that each feature carries unique information that could be useful for classification. Therefore, it is most likely that in PCA we will use a value of m as high as possible, possibly equal to the number of features in this way it will be possible to preserve as much of the information in the data as possible.

While in the case of Naive Bayes model and Full Covariance Gaussian Model could have similar performance because the features appear to be not strongly correlated. While in the case of the Tied model, which shares the covariance matrix between the classes, it may not have very good performance since the features are not strongly correlated.

2 Dimensionality reduction

Dimensionality reduction is useful in this context to compute a mapping from an n-dimensional feature space to an m-dimensional space; it is applied because it can compress information, remove noise and simplify classification.

2.1 PCA

PCA is an unsupervised dimensionality reduction technique that, given a centered dataset $X = \{x_1, \ldots, x_k\}$, it aims to find the subspace of R^n that allows to preserve most of the information, i.e. the directions with the highest variance. We start by calculating the covariance matrix:

$$\mathbf{C} = \frac{1}{K} \sum_i (x_i - \bar{x})(x_i - \bar{x})^T \tag{2}$$

After that we can compute the eigen-decomposition of $\mathbf{C} = \mathbf{U}\Sigma\mathbf{U}^T$, where $\mathbf{U}$ is the matrix of eigenvectors and Σ is the diagonal matrix of eigenvalues. Now we can project the data into the new subspace, it is spanned by the m columns of U corresponding to the m values of the highest eigenvalues.

$$\mathbf{y_i} = \mathbf{P^T}(x_i - \bar{x}) \tag{3}$$

P is the matrix corresponding to the m columns of U associated with the m highest eigenvalues of C. At this point to choose which is the best value of m we need to check how much total variance in the data we are able to retain using different values of m. To do this evaluation, it was decided to represent on a graph the variance of the data as m increases, i.e., we exploited that each eigenvalue corresponds to the variance along the corresponding axis, and the eigenvalues are the elements of the diagonal of the matrix Σ . The percentage will be calculated as the ratio of the sum of the first m eigenvalues to the sum of all eigenvalues.
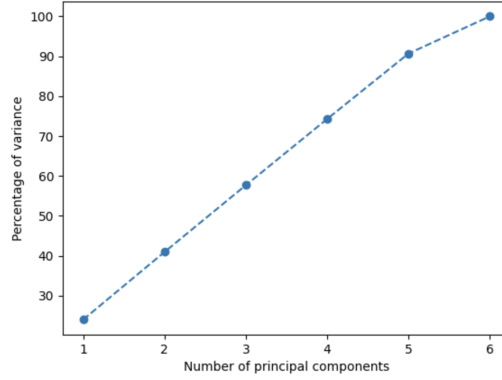

Figure 5: Variance of data for each m values

It is clear from Figure 5 that values of $m < 5$ rapidly decrease the amount of variance retained in the data. Removing a feature is probably the best choice, since this way the percentage of variance remains fairly high.
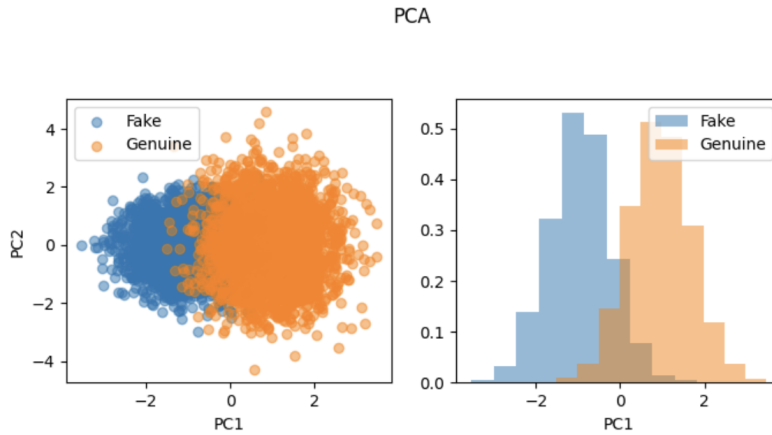

Figure 6: PCA for the two first components

After the analysis done to choose which m is best to use, we can represent the first two principal components found by calculating the PCA with m=6, visualizing these components can help to understand which features contribute most to the variance in the data. It can actually be observed that there is better separation between classes by applying PCA and choosing an m that has good variance of the data.

2.2 LDA

LDA is a supervised dimensionality reduction technique that aims to find the subspace that maximizes the separation between classes.

To compute the transformation matrix W we need to compute between-class and within-class covariance matrices.

$$\mathbf{S}_b = \frac{1}{N} \sum_{c=1}^{K} n_c (\mu_c - \mu)(\mu_c - \mu)^T \tag{4}$$

$$\mathbf{S}_w = \frac{1}{N} \sum_{c=1}^{K} \sum_{i=1}^{n_c} (x_{c,i} - \mu_i)(x_{c,i} - \mu_i)^T \tag{5}$$

where μ_c is the mean of the c-th class, μ is the mean of all the data, and n_c is the number of samples in the c-th class.

The LDA directions can be computed solving the generalized eigenvalue problem $S_b w = \lambda S_w w$, usually to solve the generalized eigenvalue problem we transform the S_w matrix into an identity matrix and then solve the eigenvalue problem for the other S_b matrix. This process allows us to find the discriminant vectors w that maximize the separation between the classes.

In LDA, the maximum number of discriminant directions or dimensions that can be obtained is always equal to C-1 where C is the number of classes. This is because the goal of the LDA is to find the directions that maximize the separation between the classes. In our case having only two classes the LDA can find at most one direction that maximizes the separation between these two classes.
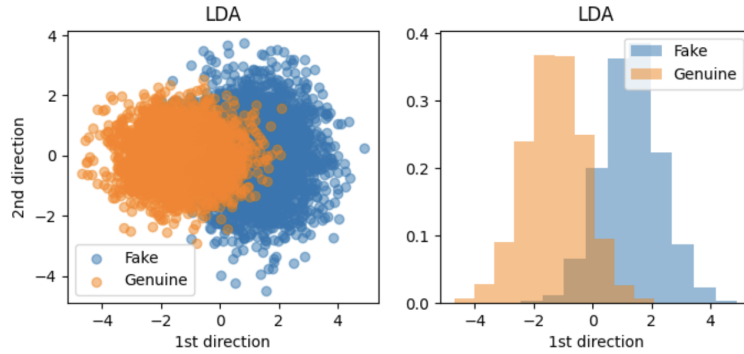

Figure 7: LDA for first direction

Looking at the results of LDA, in Figure 7, we can see that it cannot find an optimal direction that gives us a better separation of classes than the results shown previously.

2.3   Classification

Subsequent to applying dimensionality reduction techniques we used LDA to perform classification, in fact after finding the hyperplane on the training data and projecting the validation data into the new space we used different thresholds to perform classification. After applying dimensionality reduction techniques, we used LDA to perform classification; in fact, after finding the hyperplane on the training data and projecting the validation data into the new space these new data values can be used as scores by comparing them with different thresholds to perform classification, if the value exceeds the threshold it is associated with the label 1 while if it is lower 0. Two different methods were used to calculate the threshold value:

- **First Threshold**: The first method calculates the average of the LDA values for the two classes and uses the average of these two averages as the threshold. This method is simpler, but does not take into account the distribution of the data or the classification error rate. In fact, it may not provide an optimal threshold if the data are not well balanced or if the distributions of the two classes overlap significantly.

- **Second Threshold**: The second method uses a more sophisticated approach to determine the threshold. Instead of simply calculating the mean, a range of possible threshold values are examined and the one that minimizes the classification error rate is chosen. This method is more computationally heavy, but can lead to a more accurate threshold and a lower classification error rate.

Methods	Number of errors	Error rate (%)
LDA - First threshold	275	13.75
LDA - Second threshold	184	9.2
PCA(m=6)+LDA - First threshold	186	9.3
PCA(m=5)+LDA - First threshold	186	9.3
PCA(m=6)+LDA - Second threshold	184	9.2
PCA(m=5)+LDA - Second threshold	185	9.25

Table 1: LDA and PCA+LDA results with different threshold

As can be seen from the Table 1 there is not much real improvement by applying PCA as preprocessing before LDA and that the classification error turns out to be quite high so good classification is not performed using LDA or the combination of LDA with PCA.

3   Multivariate Gaussian Density

The multivariate Gaussian distribution is a generalization of the multidimensional Gaussian normal distribution. It is used to describe the probability of a vector of continuous random variables that can be correlated with each other.
The multivariate Gaussian distribution in d dimensions is characterized by a vector of averages μ and a covariance matrix Σ. The probability density function of the multivariate Gaussian distribution is given by:

$$\mathcal{N}(x|\mu, \Sigma) = \frac{1}{(2\pi)^{M/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right) \tag{6}$$

Very often to prevent numerical problems due to exponentiation of large numbers, it is usually recommended to work on the logarithm of the density:

$$\log \mathcal{N}(x|\mu, \Sigma) = -\frac{M}{2}\log(2\pi) - \frac{1}{2}\log|\Sigma| - \frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu) \tag{7}$$
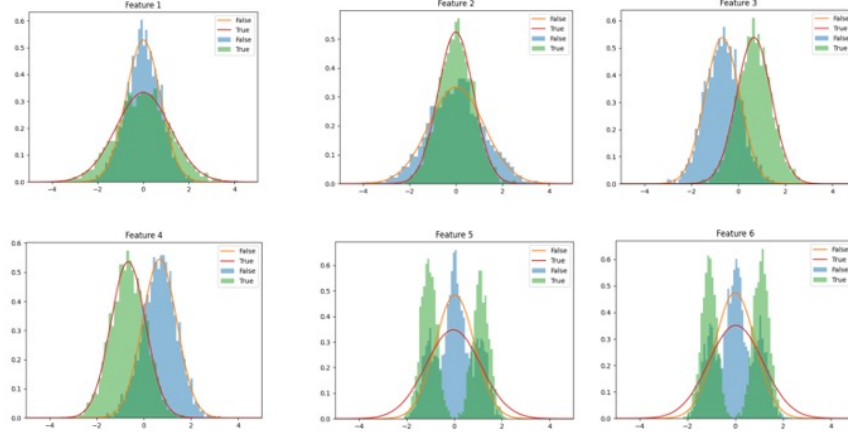

Figure 8: Gaussian density

The curve we can observe on the graph in Figure 8 represents the Gaussian probability density calculated from the data. If the data actually follow a Gaussian distribution, then the curve should fit the histogram of the data well. When this happens we can say that the Gaussian model is a good fit for the data while if the Gaussian curve does not fit the histogram well, then the Gaussian model may not be the most appropriate model for that data.

We can observe that feature 1, feature 2, feature 3 and feature 4 fit the model well and follow a Gaussian distribution.

4 Model evaluation for classification

To understand which model is the most promising and gives us the best results, we divided the dataset as mentioned above into a part for training and a part for validation. This way we can test our model on data other than the training data.

We considered different types of applications i.e. various values of $(\tilde{\pi}, C_{fn}, C_{fp})$. The target application chosen to be balanced will be:

$$(\tilde{\pi}, C_f p, C_f n) = (0.1, 1, 1) \tag{8}$$

The objective of the following analysis is to choose the model with the best performance. Performance will be measured in terms of normalized minimum detection costs. This technique would allow us to discriminate the best performing model for our target application by providing an assessment of how well the model can detect the class to which it belongs.

It is defined as follows:

$$DCF_u = \sum_{k=0}^{K-1} \frac{\pi_k}{N_k} \sum_{i|c_i=k} C(c_i^*|k) = \pi_T C_{fn} P_{fn} + (1 - \pi_T) C_{fp} P_{fp} \tag{9}$$

where $P_{fn} = \frac{FN}{FN+FP}$ and $P_{fp} = \frac{FP}{FP+TP}$.

From the Equation 4 it is possible to find **minDCF** and **actDCF**, and different thresholds are used to calculate these. In fact in the case of minDCF one uses as a threshold the one that minimizes DCF and there are several methods to be able to find it. Whereas for actDCF you use as a threshold:

$$t' = -\log \frac{\tilde{\pi}}{1 - \tilde{\pi}} \tag{10}$$

5 Classification model analysis

To perform the classification obviously we had to divide the training dataset into a part for training and a part for validation.

5.1 Gaussian Models

The first class of models we will analyze is that of Gaussian generative models. A simple model is to assume that our data, given a class, can be described by a Gaussian distribution:

$$f_{x|C}(x|c) = \mathcal{N}(x|\mu_c, \Sigma_c) \tag{11}$$

Since this is a binary classification task, we will assign a probability score to each sample in terms of the log-posterior class likelihood ratio:

$$llr(x_t) = \log\ r(x_t) = \log \frac{P(C = h_1|x_t)}{P(C = h_0|x_t)} \tag{12}$$

This expression can be expanded writing:

$$llr(x_t) = \log \frac{f_{X|C}(x_t|C = h_1)}{f_{X|C}(x_t|C = h_1)} + \frac{\pi}{1 - \pi} \tag{13}$$

where $f_{X|C}(x_t|c) = \mathcal{N}(x|\mu_c, \Sigma_c)$

It is necessary to find the parameters θ, μ_c and Σ_c; this can be done by maximizing the log-likelihood.
Parameter estimation is part of the training phase and is therefore performed on the training part of the dataset. While the estimation phase is to calculate the likelihood ratio for each sample and compare it to a threshold, which will be set at 0(th=0) in this step. After doing this, the error rate can be calculated to evaluate the classification. This analysis has been done on different Gaussian models, the difference between these models is the way the parameters (μ_c, Σ_c) are calculated, more specifically Σ_c.

MLPR Project:
Fingerprint spoofing detection

5.1.1 MVG Gaussian Classifier

The first model we will analyze is the multivariate Gaussian classifier; in this model the covariance matrix is assumed to be different for each class. The parameters μ_c and Σ_c are estimated by maximizing the log-likelihood of the training data. Thus they can be obtained as:

$$\mu_c^* = \frac{1}{N_c} \sum_{i|c_i=C} x_i \quad \sigma_c^* = \frac{1}{N_c} \sum_{i|c_i=C} (x_i - \mu_c)(x_i - \mu_c)^T \tag{14}$$

5.1.2 Naive Bayes Classifier

The Naive Bayes hypothesis simplifies the MVG full covariance model. In fact, if we know that the different components of our characteristics are approximately independent, we can simplify the estimation by assuming that the distribution of x—C can be factorized on its components.
The parameters then become:

$$\mu_{c,[j]}^* = \frac{1}{N_c} \sum_{i|c_i=C} x_{i,[j]} \quad \sigma_{c,[j]}^* = \frac{1}{N_c} \sum_{i|c_i=C} (x_{i,[j]} - \mu_{c,[j]})^2 \tag{15}$$

The Naïve Bayes classifier corresponds to the MVG full covariance classifier with a diagonal covariance matrix.

5.1.3 Tied Bayes Classifier

Another common Gaussian Model is Tied Model, it assumes that the covariance matrix is the same for all classes, but each class still has its own mean. It assumes that:

$$f_{X|C}(x|c) = \mathcal{N}(x|\mu_c, \Sigma) \tag{16}$$

While the parameters are estimated as:

$$\mu_c^* = \frac{1}{N_c} \sum_{i|c_i=C} x_i \quad \Sigma^* = \frac{1}{N} \sum_c \sum_{i|c_i=C} (x_i - \mu_c)(x_i - \mu_c)^T \tag{17}$$

where $N = \sum_{c=1}^{k} Nc$, that is the number of all samples.
This model is strongly related to LDA, if considering the binary log-likelihood ratio of the tied model we obtain a linear decision function:

$$llr(x_t) = log\frac{f_{X|C}(x|h_1)}{f_{X|C}(x|h_0)} = x^T b + c \tag{18}$$

where b and c are functions of class means and covariance matrix. On the other hand, projecting over the LDA subspace is, up to a scaling factor k, given by:

$$w^t x = k x^T \Lambda(\mu_1 - \mu_0) \tag{19}$$

where $\Lambda(\mu_1 - \mu_0) = b$. The LDA assumption that all the classes have the same within class covariance matrix is related to the assumption done for the tied model.

5.1.4 Gaussian Models Comparison

At this point placing the $P(C = 1) = P(C = 0) = 1/2$, classification was performed as can be seen in Table 2 using Gaussian models.

Features	Model	Error Rate (%)
	no PCA	
1,2,3,4,5,6	Full Cov	7.00
1,2,3,4,5,6	Naïve Bayes	7.20
1,2,3,4,5,6	Tied Cov	9.30
1,2,3,4	Full Cov	7.95
1,2,3,4	Naïve Bayes	7.65
1,2,3,4	Tied Cov	9.50
1,2	Full Cov	36.50
1,2	Naïve Bayes	36.30
1,2	Tied	49.45
3,4	Full Cov	9.45
3,4	Naïve Bayes	9.45
3,4	Tied Cov	9.40
	PCA m= 5	
1,2,3,4,5,6	Full Cov	7.10
1,2,3,4,5,6	Naïve Bayes	8.75
1,2,3,4,5,6	Tied Cov	9.30
	PCA m= 6	
1,2,3,4,5,6	Full Cov	7.00
1,2,3,4,5,6	Naïve Bayes	8.90
1,2,3,4,5,6	Tied Cov	9.30

Table 2: Error Rates for Different Models and Features

Comparing the results with Table 1 , where LDA had been used for classification we can observe that for certain configurations there was an improvement. This could be due to the fact that this method is able to classify better and thus recognize the point to which class it belongs. Analyzing the data provided in Table 2 we can make some considerations about the features used for classification:

- **Features (1,2,3,4,5,6)**: the classification error is relatively low for all models, this may suggest that all features provide useful information.

- **Features (1,2,3,4)**: the classification error increases slightly, so it can be said that features 5 and 6 probably contain useful classification information.

- **Features (1,2),(3,4)**: the classification error increases significantly in the case of features(1,2) because they do not contain particularly important information for classification. While for (3,4),the error remains similar to the others , this might indicate that these two features are particularly informative for classification.

Now we apply to estimate performance the metrics described in the previous section, we began by testing several applications to understand how changing parameters were reflected in perfor-

mance. The prior value represents the a priori probability of the positive class, so if it is higher the value is expected to be more common, while a higher value of C_{fn} indicates that the cost of misclassifying a positive point as negative is higher, and vice versa for C_{fp}.

Model	Full Cov	Naïve Bayes	Tied Cov
Application(π_T, C_{fn}, C_{fp}) : (0.50, 1, 1)			
actDCF	0.139929	0.143929	0.186044
minDCF	0.133016	0.131128	0.181244
Application(π_T, C_{fn}, C_{fp}) : (0.90, 1, 1)			
actDCF	0.400058	0.389257	0.462558
minDFC	0.342309	0.350950	0.4421083
Application(π_T, C_{fn}, C_{fp}) : (0.10, 1, 1)			
actDCF	0.304147	0.302163	0.405066
minDFC	0.262913	0.256960	0.364823
Application(π_T, C_{fn}, C_{fp}) : (0.50, 1, 9)			
actDCF	0.305140	0.302163	0.406058
minDFC	0.262913	0.256960	0.362839
Application(π_T, C_{fn}, C_{fp}) : (0.50, 9, 1)			
actDCF	0.400058	0.389257	0.462558
minDFC	0.342309	0.350950	0.445132

Table 3: minDCF and actDCF for Different Models and configurations

So it can be seen from the Table 3 that:

- When the value of Prior increases (from 0.50 to 0.90), a worsening of the models can be seen. This is due to the fact that the model is penalizing false negatives more heavily. This also happens when Prior decreases(from 0.50 to 0.10) and therefore you are penalizing false positives more heavily.

- When C_{fp} or C_{fn} increase to 9.00, the actDCF and minDCF of all models worsen compared with the case when C_{fp} and C_{fn} were 1.00. However, it is possible to see the increase in the cost of false positives has less significant impact on the performance of the model so this suggests that the increase in the cost of false positives affects the model less. While in the case of the cost of false negatives this has a higher impact so from this it can be inferred that increasing the cost of false negatives worsens the performance of the models.

We now focus on the three applications that have different effective priors (0.1, 0.5, 0.9) and with costs of errors equal to 1.00 and we can, we can consider PCA as pre-processing.

Model	Full Cov	Naïve Bayes	Tied Cov
Application(π_T, C_{fn}, C_{fp}) : (0.50, 1, 1)			
no PCA			
actDCF	0.139929	0.143929	0.186044
minDCF	0.133016	0.131128	0.181244
PCA			
m=5			
actDCF	0.141913	0.174987	0.186044
minDCF	0.133144	0.173691	0.181163
m=6			
actDCF	0.139929	0.177995	0.186028
minDCF	0.130168	0.172699	0.181244
Application(π_T, C_{fn}, C_{fp}) : (0.90, 1, 1)			
no PCA			
actDCF	0.400058	0.389257	0.462558
minDFC	0.342309	0.350950	0.4421083
PCA			
m=5			
actDCF	0.398041	0.466014	0.462558
minDCF	0.351238	0.434044	0.445132
m=6			
actDCF	0.400058	0.451181	0.462558
minDCF	0.342309	0.435915	0.442108
Application(π_T, C_{fn}, C_{fp}) : (0.10, 1, 1)			
no PCA			
actDCF	0.304147	0.302163	0.405066
minDFC	0.262913	0.256960	0.364823
PCA			
m=5			
actDCF	0.304147	0.393017	0.405066
minDFC	0.273825	0.354470	0.364823
m=6			
actDCF	0.305140	0.392025	0.406058
minDCF	0.262913	0.353479	0.362839

Table 4: minDCF and actDCF for Different Models and configurations with and without PCA

- From this first analysis, we can infer that the worst model is Tied Covariance (Tied Cov), as we had already assumed by analyzing the correlation of features since in this model the features are assumed to be correlated and in our case they are not strongly correlated.

- We can see that after applying PCA, actDCF and minDCF remain unchanged or in some cases even worsen, which could indicate that the reduction in dimensionality removed some important information from the data that was useful for classification. That is, the components (or dimensions) that have been removed might contain discriminative information that

helps the model distinguish between classes, or the transformation might break the assumptions made for the models under analysis and thus we might get an inaccurate estimate that worsens performance.

- According to different applications the full Covariance and naive bayes model have similar performance, however, both have some strengths and weaknesses. The Full Covariance model considers all correlations between features, so it may be able to capture more complex relationships between features, but this may lead to overfitting and also may be computationally heavier. On the other hand, on the other hand, the Naive Bayes model assumes that all features are independent of each other which might not be entirely true in our situation but it gets us good performance and might be a simplification, this however is less computationally heavy.

- The application that has the best calibration, that is, where minDCF and actDCF are as close as possible is the one with the $\pi_T = 0.5$ both in the case where there is no pca application and the one where there is. In this case we can say that there is a good calibration for almost all models, while in the case of the other priors there is no good calibration.
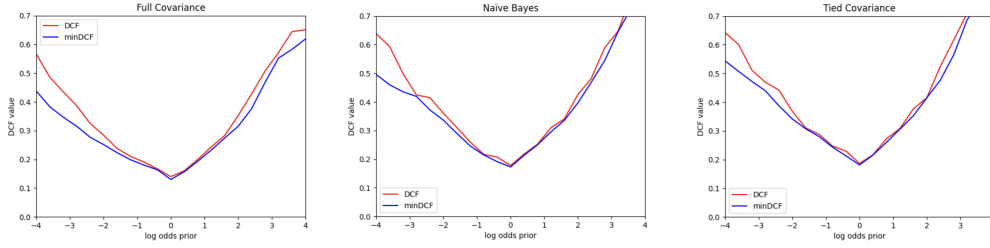

Figure 9: Confronto tra i modelli: Full Covariance, Naive Bayes, e Tied Covariance

<u>Best Gaussian Model</u> analyzed for our main application that is application with $\pi_T = 0.1$ is:

- **PCA**: Model Full Covariance with PCA(m=6)

- **No PCA**: Model Naive Bayes

5.2   Logistic Regression Classifier

Logistic Regression is a discriminative classification model that makes no assumptions about the distribution of the data, but tries to find a hyperplane that maximizes the posterior likelihood. Starting from the results obtained from the Tied Gaussian that provides log-likelihood ratios that are linear functions of our data we can write it as:

$$llr(x) = \log \frac{P(C = h_1|x)}{P(C = h_0|x)} = \log \frac{f_{X|C}(x|h_1)}{f_{X|C}(x|h_0)} + \log \frac{\pi}{1 - \pi} = w^T x + b \tag{20}$$

All prior information has been absorbed in the bias term b, given w and b we can compute the expression for the posterior class probability:

$$P(C = h_1|x, w, b) = \frac{e^{(w^T x + b)}}{1 + e^{(w^T x - b)}} = \sigma(w^T x + b) \tag{21}$$

where $\sigma = \frac{1}{1-e^{-x}}$ is the sigmoid function.

This equation provides a model that allows computing the posterior prabilities for h1 and h0, the model assumes that decision rules are linear surfaces orthogonal to the vector **w** and the model parameters are $(\mathbf{w}, b)$.

5.2.1 Binary Logistic Regression

Binary Logistic Regression non prior-weighted

We are going to look for the minimizer of the function:

$$\mathbf{J}(w, b) = \frac{\lambda}{2}||w||^2 + \frac{1}{n}\sum_{i=1}^{n}\log(1 + e^{-z_i(w^T x_i + b)}), \quad z_i = \begin{cases} 1 & \text{if } c_i = 1 \\ -1 & \text{if } c_i = 0 \end{cases} \tag{22}$$

where λ is an hyperparameter that represents the regularization term (or norm penalty), that has been introduced to make the problem solvable in case of linearly separable classes.
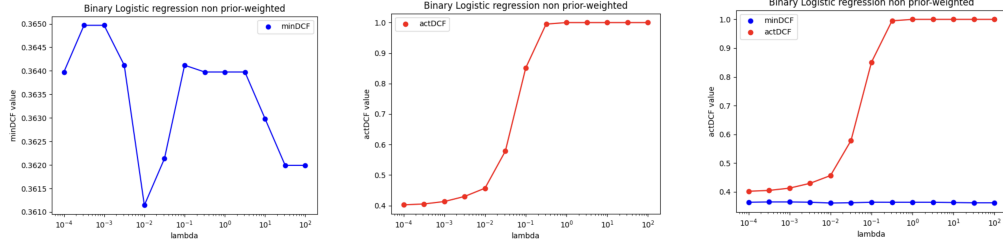

Figure 10: Binary Logistic Regression Model non prior-weighted
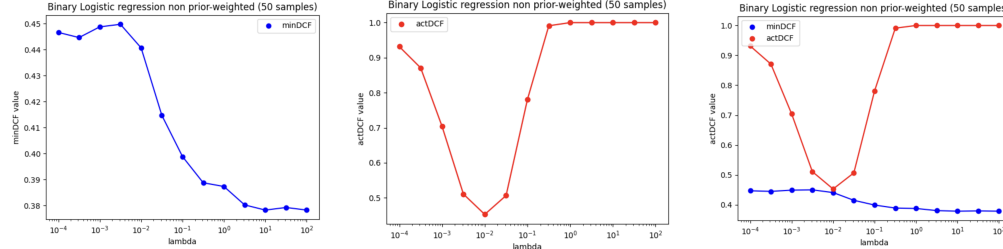

Figure 11: Binary Logistic Regression Model non prior-weighted(50 samples)

Now we can analyze the results of the binary logistic regression model on the primary application $\pi_T = 0.1$.

λ	minDCF	actDCF
Binary Logistic Regression non prior-weighted model		
10^{-4}	0.363975	0.402089
10^{-3}	0.364967	0.413002
10^{-2}	0.361143	0.413002
10^{-1}	0.364119	0.851190
Binary Logistic Regression non prior-weighted (50 samples)		
10^{-4}	0.446604	0.93146
10^{-3}	0.448733	0.704077
10^{-2}	0.440652	0.452557
10^{-1}	0.398842	0.780898

Table 5: minDCF and actDCF

- Using only 50 samples, we can see that using a limited number of samples can significantly affect the logistic regression model and could lead to misleading outcomes that are not representative of the entire dataset. In fact, the λ values for which we get better performance turn out to be different from those obtained in the table 5 in the section 'Binary Logistic Regression non prior-weighted model' where we used the entire dataset. From this we can conclude that there is a need to use as many samples as possible for training in order to obtain a more accurate model.

Binary Logistic Regression with prior-weighted

It's also possible use prior-weighted logistic regression model that allows us to simulate different priors for class 1. So the objective function becomes:

$$\mathbf{J}(w,b) = \frac{\lambda}{2}||w||^2 + \frac{1}{n}\sum_{i=1}^{n}\xi_i \log(1 + e^{-z_i(w^T x_i + b)}), \quad \xi_i = \begin{cases} \frac{\pi_T}{n_T} & \text{if } z_i = +1(c_i = 1) \\ \frac{1-\pi_T}{n_F} & \text{if } z_i = -1(c_i = 0) \end{cases} \quad (23)$$
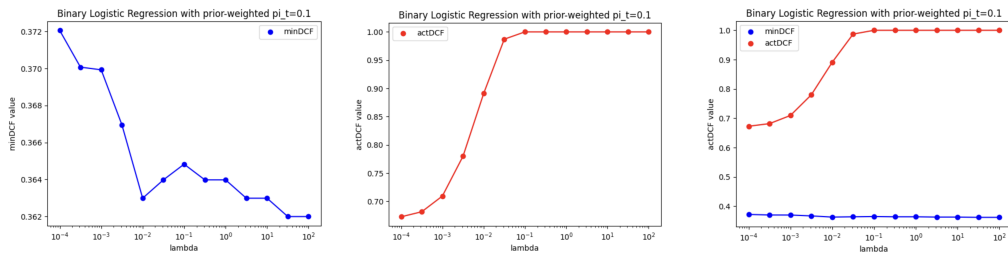

Figure 12: Binary Logistic Regression Model with prior-weighted prior=0.1

Now we can analyze the results of the binary logistic regression model on the primary application $\pi_T = 0.1$.

λ	minDCF	actDCF
Binary Logistic Regression with prior-weighted		
$\pi_T = 0.1$		
10^{-4}	0.372056	0.672763
10^{-3}	0.369928	0.709469
10^{-2}	0.362983	0.890873
10^{-1}	0.364823	1.0

Table 6: minDCF and actDCF

- In logistic regression, prior is a parameter that represents our a priori knowledge about the distribution of labels in the dataset. In other words, it is the probability that a random sample belongs to a certain class before we have seen the data. So in our current case we use the prior to weight the samples during model training. In particular, samples in the higher priority class receive a higher weight than those in the lower priority class. This can be useful if the dataset is unbalanced, the choice of the prior must be made at the beginning and this choice can affect the model a lot, in fact we may even have a worsening of the model.

- In our case for example we do not see a significant model worsening, in fact our dataset is not particularly unbalanced, so we can see a significant worsening of the actDCF by applying binary logistic regression with prior-weighted $\pi_T = 0.1$.

Binary Logistic Regression with Pre-Processing (PCA)

We can now analyze the results obtained by applying the model after applying PCA, always using $\pi_T = 0.1$

λ	minDCF	actDCF
Binary Logistic Regression with PCA		
m=5		
10^{-4}	0.366103	0.401098
10^{-3}	0.366103	0.410026
10^{-2}	0.361847	0.457773
10^{-1}	0.365959	0.849206
m=6		
10^{-4}	0.363975	0.402089
10^{-3}	0.364967	0.413002
10^{-2}	0.361143	0.456781
10^{-1}	0.364119	0.851190

Table 7: minDCF and actDCF

We can conclude by looking at the table that the application of PCA does not give a significant improvement, as noted earlier, but we can state that performance remains constant.

5.2.2 Quadratic Logistic Regression

Now we can analyze training on a Quadratic Logistic Regression model by performing features expansion. We can write log-likelihood ratio as:

$$\log \frac{P(C = h_1|x)}{P(C = h_0|x)} = x^T A x + b^T x + c = s(\mathbf{x}, \mathbf{A}, \mathbf{b}, \mathbf{c}) \tag{24}$$

This expression is quadratic in x but it's linear in A and b. We could rewrite it to obtain a decision function that is linear for the expanded features space but quadratic in original features space. We can write features expansion as:

$$\Phi(x) = \begin{bmatrix} vec(xx^T) \\ x \end{bmatrix}, \quad w = \begin{bmatrix} vec(A) \\ b \end{bmatrix} \tag{25}$$

where vec(X) is the operator that stacks the columns of X into a single column vector. In this way we can write the posterio log-likelihood as:
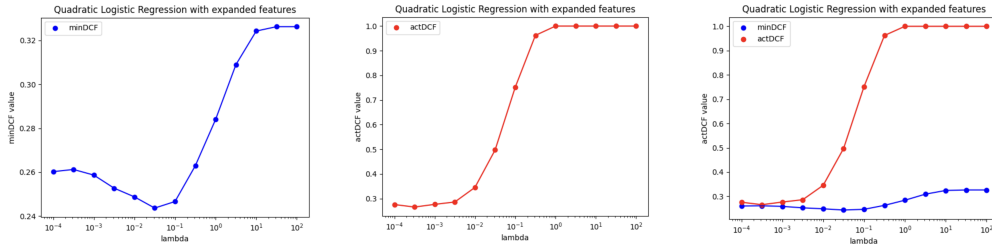
$$s(x, w, c) = s^T \phi(x) + c \tag{26}$$

Figure 13: Quadratic Logistic Regression

λ	minDCF	actDCF
10^{-4}	0.260224	0.275809
10^{-3}	0.258656	0.276514
10^{-2}	0.248736	0.345382
10^{-1}	0.246607	0.751984

Table 8: minDCF and actDCF

The quadratic model can help us extract correlations between features that cannot be extracted by a linear model, in our case we can see that Quadratic Logistic Regression performs better than Binary Logistic Regression with and without prior-weighted.

Conclusion for Binary Logistic Regression

- For all the various models represented in Figure 10 , Figure 12 , Figure 13 , it can be seen that λ significantly affects actDCF and minDCF. In fact, it can be seen that for values greater than 10^{-1} there is significant degradation. It is necessary to remember that a larger value of λ may lead to overgeneralization which may lead to underfitting, while for too small

values of it there may be low generalization and thus there may be overfitting. Based on this analysis, large values of λ were excluded in the tables showing the different values of minDCF and actDCF for lambda ranging from 10^{-4} to 10^{-1}.

- After doing various analyses we can see which are the optimal values of lambda as mentioned above and which model gives the best performance comparing for now the minDCF, also the difference we can see between minDCF and actDCF suggests i that the models do not have a good calibration. In our case quadratic logistic regression seems to give the best results, so we can draw that we need the model to be able to capture nonlinear relationships between variables and thus be able to find nonlinear separations rules in the feature space.

- From the results we can observe that in fact we get better results with a relatively higher lambda so we can conclude that we need more regularization to prevent overfitting.

Logistic Regression Model analyzed for our main application that is application with $\pi_T = 0.1$ is:

- **Quadratic Logistic Regression** with $\lambda = 10^{-1}$, no prior-weighted

5.3 SVM Classifier

5.3.1 Linear SVM

5.3.2 Kernel SVM

5.4 GMM Classifier

6 Score Calibration

6.1 Calibration Analysis on Selected Models

6.2 Calibrating Scores for Selected Models

7 Experimental Results

7.1 Calibration on evaluation score

7.2 Considerations

8 Conclusions